

```
In [2]: import numpy as np
from keras.datasets import cifar10
from keras.models import Sequential
from keras.layers.core import Dense, Flatten
from keras.layers.convolutional import Conv2D
from tensorflow.keras.optimizers import Adam
from keras.layers.pooling import MaxPooling2D
from tensorflow.keras.utils import to_categorical
import matplotlib.pyplot as plt

In [3]: np.random.seed(42)
import ssl
ssl._create_default_https_context = ssl._create_unverified_context


In [4]: (X_train, Y_train), (X_test, Y_test) = cifar10.load_data()

Downloading data from https://www.cs.toronto.edu/~kriz/cifar-10-python.tar.gz
17059896747048971 [=====] - 925.1us/step
17059896747048971 [=====] - 925.1us/step

In [4]: class_names = ['airplane', 'automobile', 'bird', 'cat', 'deer',
                        'dog', 'frog', 'horse', 'ship', 'truck']

fig = plt.figure(figsize=(8, 3))
for i in range(len(class_names)):
    ix = fig.add_subplot(2, 5, 1 + i, ticks=[])
    idx = np.where(Y_train == i)[0]
    features_idx = X_train[idx,:]
    img_num = np.random.randint(features_idx.shape[0])
    ix = features_idx[idx_num,:]
    ix.set_title(class_names[i])
    img = np.transpose(features_idx[idx_num,:], (1, 2, 0))
    plt.imshow(img)

plt.show()
```



Initializing model with no normalization

```
In [9]: # Initializing the model
model = Sequential()
# defining a convolutional layer
model.add(Conv2D(32, kernel_size=(3, 3), activation='relu', input_shape=(32, 32, 3)))
# defining a second convolutional layer
model.add(Conv2D(32, kernel_size=(3, 3), activation='relu'))
# defining a third convolutional layer
model.add(Conv2D(32, kernel_size=(3, 3), activation='relu'))
# we add our classifier
model.add(Flatten())
model.add(Dense(1024, activation='relu'))
model.add(Dense(10, activation='softmax'))
# compiling the model
model.compile(loss='categorical_crossentropy',
              optimizer=Adam(learning_rate=0.001, decay=1e-6),
              metrics=['accuracy'])

# Training of the model
model.fit(X_train, to_categorical(Y_train),
        batch_size=128,
        shuffle=True,
        epochs=5,
        validation_data=(X_test, to_categorical(Y_test)))

# Evaluation of the model
scores = model.evaluate(X_test, to_categorical(Y_test))
print('Loss: %.3f' % scores[0])
print('Accuracy: %.3f' % scores[1])

Epoch 1/5
391/391 [=====] - 10795.3s/step - loss: 7.4783 - accuracy: 0.4685 - val_loss: 1.2216 - val_accuracy: 0.5600
Epoch 2/5
391/391 [=====] - 6095.2s/step - loss: 0.9259 - accuracy: 0.6808 - val_loss: 1.0981 - val_accuracy: 0.6252
Epoch 3/5
391/391 [=====] - 6655.2s/step - loss: 0.5238 - accuracy: 0.8274 - val_loss: 1.1936 - val_accuracy: 0.6135
Epoch 4/5
391/391 [=====] - 9665.2s/step - loss: 0.2318 - accuracy: 0.9385 - val_loss: 1.3367 - val_accuracy: 0.6434
Epoch 5/5
391/391 [=====] - 11474.3s/step - loss: 0.0719 - accuracy: 0.9826 - val_loss: 1.5635 - val_accuracy: 0.6434
313/313 [=====] - 745.237s/step - loss: 1.5635 - accuracy: 0.6434
Loss: 1.563
Accuracy: 0.643

Data Pre-processing and normalization

Normalization is only done on the training data!
```

```
In [5]: # Centering the data
X_train_mean = np.mean(X_train, axis=0)
X_train_cent = X_train - X_train_mean
# Normalization
X_train_std = np.std(X_train, axis=0)
X_train_norm = X_train_cent / X_train_std

In [6]: X_test_norm = (X_test - X_train_mean) / X_train_std
```

Running model again with data normalization

```
In [8]: # Initializing the model
model = Sequential()
# defining a convolutional layer
model.add(Conv2D(32, kernel_size=(3, 3), activation='relu', input_shape=(32, 32, 3)))
# defining a second convolutional layer
model.add(Conv2D(32, kernel_size=(3, 3), activation='relu'))
# defining a third convolutional layer
model.add(Conv2D(32, kernel_size=(3, 3), activation='relu'))
# we add our classifier
model.add(Flatten())
model.add(Dense(1024, activation='relu'))
model.add(Dense(10, activation='softmax'))
# compiling the model
model.compile(loss='categorical_crossentropy',
              optimizer=Adam(learning_rate=0.001, decay=1e-6),
              metrics=['accuracy'])

# Training of the model
model.fit(X_train_norm, to_categorical(Y_train),
        batch_size=128,
        shuffle=True,
        epochs=5,
        validation_data=(X_test, to_categorical(Y_test)))

# Evaluation of the model
scores = model.evaluate(X_test, to_categorical(Y_test))
print('Loss: %.3f' % scores[0])
print('Accuracy: %.3f' % scores[1])

Epoch 1/5
391/391 [=====] - 4685.1s/step - loss: 5.7654 - accuracy: 0.4794 - val_loss: 1.1737 - val_accuracy: 0.5999
Epoch 2/5
391/391 [=====] - 4645.1s/step - loss: 0.8229 - accuracy: 0.7225 - val_loss: 1.1932 - val_accuracy: 0.6372
Epoch 3/5
391/391 [=====] - 4595.1s/step - loss: 0.3568 - accuracy: 0.8889 - val_loss: 1.2649 - val_accuracy: 0.6438
Epoch 4/5
391/391 [=====] - 4755.1s/step - loss: 0.1098 - accuracy: 0.9710 - val_loss: 1.4464 - val_accuracy: 0.6329
Epoch 5/5
391/391 [=====] - 4566.1s/step - loss: 0.0352 - accuracy: 0.9932 - val_loss: 1.6619 - val_accuracy: 0.6337
313/313 [=====] - 139.686s/step - loss: 1.6619 - accuracy: 0.6337
Loss: 1.662
Accuracy: 0.634

I can see that this helped make the training accuracy better but the validation accuracy is still low. We will now try different regularization techniques to make it less overfit
```

Batch Normalization

```
In [10]: # We import Batch Normalization layer
from tensorflow.keras.layers import BatchNormalization, Activation
# Initializing the model
model = Sequential()
# defining a convolutional layer
model.add(Conv2D(32, kernel_size=(3, 3), input_shape=(32, 32, 3)))
model.add(BatchNormalization())
# defining a second convolutional layer
model.add(Conv2D(32, kernel_size=(3, 3), activation='relu'))
model.add(BatchNormalization())
# defining a third convolutional layer
model.add(Conv2D(32, kernel_size=(3, 3), activation='relu'))
model.add(BatchNormalization())
# we include our classifier
model.add(Flatten())
model.add(Dense(1024, activation='relu'))
model.add(Dense(10, activation='softmax'))
# compiling the model
model.compile(loss='categorical_crossentropy',
              optimizer=Adam(lr=0.001, decay=1e-6),
              metrics=['accuracy'])

# Training the model
model.fit(X_train_norm, to_categorical(Y_train),
        batch_size=128,
        shuffle=True,
        epochs=5,
        validation_data=(X_test_norm, to_categorical(Y_test))) # aqui deberíamos usar un conjunto distinto al de test!!

# Evaluating the model
scores = model.evaluate(X_test_norm, to_categorical(Y_test))
print('Loss: %.3f' % scores[0])
print('Accuracy: %.3f' % scores[1])

Epoch 1/5
391/391 [=====] - 5825.1s/step - loss: 1.5639 - accuracy: 0.5243 - val_loss: 1.4934 - val_accuracy: 0.5286
Epoch 2/5
391/391 [=====] - 5795.1s/step - loss: 0.7281 - accuracy: 0.7501 - val_loss: 0.8749 - val_accuracy: 0.6535
Epoch 3/5
391/391 [=====] - 5755.1s/step - loss: 0.3378 - accuracy: 0.9073 - val_loss: 0.8740 - val_accuracy: 0.6763
Epoch 4/5
391/391 [=====] - 5935.2s/step - loss: 0.1075 - accuracy: 0.9859 - val_loss: 0.8101 - val_accuracy: 0.6749
Epoch 5/5
391/391 [=====] - 6055.2s/step - loss: 0.0387 - accuracy: 0.9902 - val_loss: 1.0356 - val_accuracy: 0.6964
313/313 [=====] - 265.838s/step - loss: 1.0356 - accuracy: 0.6964
Loss: 1.036
Accuracy: 0.696

I can see that our training accuracy is pretty much perfect and that our validation accuracy improved slightly but can still get better. Using this also did make the neural network get trained faster.
```

Lasso Regularization (L1)

```
In [12]: # L2 Regularization
# Regularizer Layer import
from keras.regularizers import l2
# Initializing the model
model = Sequential()
# defining a convolutional layer
model.add(Conv2D(32, kernel_size=(3, 3), activation='relu', input_shape=(32, 32, 3)))
# defining a second convolutional layer
model.add(Conv2D(32, kernel_size=(3, 3), activation='relu'))
# defining a third convolutional layer
model.add(Conv2D(32, kernel_size=(3, 3), activation='relu'))
# Classifier inclusion
model.add(Flatten())
model.add(Dense(1024, activation='relu', kernel_regularizer=l2(0.01)))
model.add(Dense(10, activation='softmax'))
# compiling the model
model.compile(loss='categorical_crossentropy',
              optimizer=Adam(lr=0.001, decay=1e-6),
              metrics=['accuracy'])

# Training the model
model.fit(X_train_norm, to_categorical(Y_train),
        batch_size=128,
        shuffle=True,
        epochs=5,
        validation_data=(X_test_norm, to_categorical(Y_test)))

# Evaluating the model
scores = model.evaluate(X_test_norm, to_categorical(Y_test))
print('Loss: %.3f' % scores[0])
print('Accuracy: %.3f' % scores[1])

Epoch 1/5
391/391 [=====] - 6385.2s/step - loss: 3.6848 - accuracy: 0.4655 - val_loss: 1.4961 - val_accuracy: 0.5443
Epoch 2/5
391/391 [=====] - 6215.2s/step - loss: 1.4393 - accuracy: 0.5723 - val_loss: 1.3997 - val_accuracy: 0.5719
Epoch 3/5
391/391 [=====] - 6185.2s/step - loss: 1.3188 - accuracy: 0.6135 - val_loss: 1.3195 - val_accuracy: 0.5986
Epoch 4/5
391/391 [=====] - 6165.2s/step - loss: 1.2477 - accuracy: 0.6356 - val_loss: 1.2346 - val_accuracy: 0.6471
Epoch 5/5
391/391 [=====] - 6105.2s/step - loss: 1.2028 - accuracy: 0.6544 - val_loss: 1.1838 - val_accuracy: 0.6625
313/313 [=====] - 555.174s/step - loss: 1.1839 - accuracy: 0.6625
Loss: 1.184
Accuracy: 0.663

This did not really help since the accuracy score is actually a little lower and the loss amount is similar
```

Ridge Regularization (L2)

```
In [13]: # L2 Regularization
# Regularizer Layer import
from keras.regularizers import l1
# Initializing the model
model = Sequential()
# defining a convolutional layer
model.add(Conv2D(32, kernel_size=(3, 3), activation='relu', input_shape=(32, 32, 3)))
# defining a second convolutional layer
model.add(Conv2D(32, kernel_size=(3, 3), activation='relu'))
# defining a third convolutional layer
model.add(Conv2D(32, kernel_size=(3, 3), activation='relu'))
# Classifier inclusion
model.add(Flatten())
model.add(Dense(1024, activation='relu', kernel_regularizer=l1(0.01)))
model.add(Dense(10, activation='softmax'))
# compiling the model
model.compile(loss='categorical_crossentropy',
              optimizer=Adam(lr=0.001, decay=1e-6),
              metrics=['accuracy'])

# Training the model
model.fit(X_train_norm, to_categorical(Y_train),
        batch_size=128,
        shuffle=True,
        epochs=5,
        validation_data=(X_test_norm, to_categorical(Y_test)))

# Evaluating the model
scores = model.evaluate(X_test_norm, to_categorical(Y_test))
print('Loss: %.3f' % scores[0])
print('Accuracy: %.3f' % scores[1])

Epoch 1/5
391/391 [=====] - 6215.2s/step - loss: 284.0510 - accuracy: 0.2429 - val_loss: 12.8222 - val_accuracy: 0.2622
Epoch 2/5
391/391 [=====] - 6205.2s/step - loss: 12.5892 - accuracy: 0.3172 - val_loss: 12.3606 - val_accuracy: 0.3493
Epoch 3/5
391/391 [=====] - 6215.2s/step - loss: 12.1354 - accuracy: 0.3581 - val_loss: 12.0529 - val_accuracy: 0.3890
Epoch 4/5
391/391 [=====] - 6225.2s/step - loss: 12.0569 - accuracy: 0.3813 - val_loss: 12.0312 - val_accuracy: 0.3833
Epoch 5/5
391/391 [=====] - 6215.2s/step - loss: 11.8643 - accuracy: 0.4042 - val_loss: 11.7613 - val_accuracy: 0.4099
313/313 [=====] - 555.170s/step - loss: 11.7613 - accuracy: 0.4099
Loss: 11.760
Accuracy: 0.409

This had extremely low accuracy and a high loss when compared to the other regularization methods
```

Elastic Net Regularization (L1+L2)

```
In [14]: # Elastic Net Regularization (L1 + L2)
# Regularizer Layer import
from keras.regularizers import l1_l2
# Initializing the model
model = Sequential()
# defining a convolutional layer
model.add(Conv2D(32, kernel_size=(3, 3), activation='relu', input_shape=(32, 32, 3)))
# defining a second convolutional layer
model.add(Conv2D(32, kernel_size=(3, 3), activation='relu'))
# defining a third convolutional layer
model.add(Conv2D(32, kernel_size=(3, 3), activation='relu'))
# Classifier inclusion
model.add(Flatten())
model.add(Dense(1024, activation='relu', kernel_regularizer=l1_l2(0.01, 0.01)))
model.add(Dense(10, activation='softmax'))
# compiling the model
model.compile(loss='categorical_crossentropy',
              optimizer=Adam(lr=0.001, decay=1e-6),
              metrics=['accuracy'])

# Training the model
model.fit(X_train_norm, to_categorical(Y_train),
        batch_size=128,
        shuffle=True,
        epochs=5,
        validation_data=(X_test_norm, to_categorical(Y_test)))

# Evaluating the model
scores = model.evaluate(X_test_norm, to_categorical(Y_test))
print('Loss: %.3f' % scores[0])
print('Accuracy: %.3f' % scores[1])

Epoch 1/5
391/391 [=====] - 7445.2s/step - loss: 285.2887 - accuracy: 0.2386 - val_loss: 12.7639 - val_accuracy: 0.2368
Epoch 2/5
391/391 [=====] - 7405.2s/step - loss: 12.4869 - accuracy: 0.3035 - val_loss: 12.0038 - val_accuracy: 0.2906
Epoch 3/5
391/391 [=====] - 7425.2s/step - loss: 11.9379 - accuracy: 0.3432 - val_loss: 11.8223 - val_accuracy: 0.3224
Epoch 4/5
391/391 [=====] - 7375.2s/step - loss: 11.7063 - accuracy: 0.3641 - val_loss: 11.6211 - val_accuracy: 0.3830
Epoch 5/5
391/391 [=====] - 7345.2s/step - loss: 11.5752 - accuracy: 0.3779 - val_loss: 11.4863 - val_accuracy: 0.3928
313/313 [=====] - 895.283s/step - loss: 11.4863 - accuracy: 0.3928
Loss: 11.486
Accuracy: 0.393

Combining the two did not seem to help very much as the accuracy was still low with a high loss
```

Max Norm Constraints

```
In [17]: # Regularizer Layer import
from keras.constraints import max_norm
# Initializing the model
model = Sequential()
# defining a convolutional layer
model.add(Conv2D(32, kernel_size=(3, 3), activation='relu', input_shape=(32, 32, 3)))
# defining a second convolutional layer
model.add(Conv2D(32, kernel_size=(3, 3), activation='relu'))
# defining a third convolutional layer
model.add(Conv2D(32, kernel_size=(3, 3), activation='relu'))
# Classifier inclusion
model.add(Flatten())
model.add(Dense(1024, activation='relu', kernel_constraint=max_norm(3.)))
model.add(Dense(10, activation='softmax'))
# compiling the model
model.compile(loss='categorical_crossentropy',
              optimizer=Adam(lr=0.001, decay=1e-6),
              metrics=['accuracy'])

# Training the model
model.fit(X_train_norm, to_categorical(Y_train),
        batch_size=128,
        shuffle=True,
        epochs=5,
        validation_data=(X_test_norm, to_categorical(Y_test)))

# Evaluating the model
scores = model.evaluate(X_test_norm, to_categorical(Y_test))
print('Loss: %.3f' % scores[0])
print('Accuracy: %.3f' % scores[1])

Epoch 1/5
391/391 [=====] - 6395.2s/step - loss: 1.3961 - accuracy: 0.5950 - val_loss: 1.1465 - val_accuracy: 0.5994
Epoch 2/5
391/391 [=====] - 6185.2s/step - loss: 0.9810 - accuracy: 0.6610 - val_loss: 0.9853 - val_accuracy: 0.6575
Epoch 3/5
391/391 [=====] - 6245.2s/step - loss: 0.7688 - accuracy: 0.7363 - val_loss: 0.9401 - val_accuracy: 0.6755
Epoch 4/5
391/391 [=====] - 6075.2s/step - loss: 0.5863 - accuracy: 0.8007 - val_loss: 0.9161 - val_accuracy: 0.6890
Epoch 5/5
391/391 [=====] - 6125.2s/step - loss: 0.4137 - accuracy: 0.8805 - val_loss: 0.9547 - val_accuracy: 0.6882
313/313 [=====] - 295.686s/step - loss: 0.9547 - accuracy: 0.6882
Loss: 0.955
Accuracy: 0.688

Using constraints was an interesting way to perform regularization. While the accuracy was slightly lower than with just the batch normalization, the loss was slightly lower. I can also see that it is not nearly as overfit as the other methods as the accuracy and validation accuracy are not nearly as far apart
```

Dropout regularization

```
In [18]: # Dropout
# Dropout Layer import
from keras.layers import Dropout
# Initializing the model
model = Sequential()
# defining a convolutional layer
model.add(Conv2D(32, kernel_size=(3, 3), activation='relu', input_shape=(32, 32, 3)))
# defining a second convolutional layer
model.add(Conv2D(32, kernel_size=(3, 3), activation='relu'))
# defining a third convolutional layer
model.add(Conv2D(32, kernel_size=(3, 3), activation='relu'))
# Classifier inclusion
model.add(Flatten())
model.add(Dense(1024, activation='relu'))
model.add(Dropout(0.25))
model.add(Dense(10, activation='softmax'))
# compiling the model
model.compile(loss='categorical_crossentropy',
              optimizer=Adam(lr=0.001, decay=1e-6),
              metrics=['accuracy'])

# Training the model
model.fit(X_train_norm, to_categorical(Y_train),
        batch_size=128,
        shuffle=True,
        epochs=5,
        validation_data=(X_test_norm, to_categorical(Y_test)))

# Evaluating the model
scores = model.evaluate(X_test_norm, to_categorical(Y_test))
print('Loss: %.3f' % scores[0])
print('Accuracy: %.3f' % scores[1])

Epoch 1/5
391/391 [=====] - 5775.1s/step - loss: 1.5625 - accuracy: 0.4401 - val_loss: 1.2677 - val_accuracy: 0.5531
Epoch 2/5
391/391 [=====] - 5595.1s/step - loss: 1.2038 - accuracy: 0.5759 - val_loss: 1.0808 - val_accuracy: 0.6194
Epoch 3/5
391/391 [=====] - 5585.1s/step - loss: 1.0097 - accuracy: 0.6462 - val_loss: 0.9475 - val_accuracy: 0.6711
Epoch 4/5
391/391 [=====] - 5555.1s/step - loss: 0.8579 - accuracy: 0.7051 - val_loss: 0.8866 - val_accuracy: 0.6922
Epoch 5/5
391/391 [=====] - 5545.1s/step - loss: 0.7299 - accuracy: 0.7470 - val_loss: 0.8675 - val_accuracy: 0.7038
313/313 [=====] - 155.686s/step - loss: 0.8675 - accuracy: 0.7038
Loss: 0.867
Accuracy: 0.702

This was very effective as it was able to get the value accuracy to be greater than 70% and a steep drop-off in the loss
```

Max Norm + Dropout

```
In [19]: # Dropout & Max Norm
# Dropout & Max Norm layers import
from keras.layers import Dropout
from keras.constraints import max_norm
# Initializing the model
model = Sequential()
# defining a convolutional layer
model.add(Conv2D(32, kernel_size=(3, 3), activation='relu', input_shape=(32, 32, 3)))
# defining a second convolutional layer
model.add(Conv2D(32, kernel_size=(3, 3), activation='relu'))
# defining a third convolutional layer
model.add(Conv2D(32, kernel_size=(3, 3), activation='relu'))
# Classifier inclusion
model.add(Flatten())
model.add(Dense(1024, activation='relu', kernel_constraint=max_norm(3.)))
model.add(Dropout(0.5))
model.add(Dense(10, activation='softmax'))
# compiling the model
model.compile(loss='categorical_crossentropy',
              optimizer=Adam(lr=0.001, decay=1e-6),
              metrics=['accuracy'])

# Training the model
model.fit(X_train_norm, to_categorical(Y_train),
        batch_size=128,
        shuffle=True,
        epochs=5,
        validation_data=(X_test_norm, to_categorical(Y_test)))

# Evaluating the model
scores = model.evaluate(X_test_norm, to_categorical(Y_test))
print('Loss: %.3f' % scores[0])
print('Accuracy: %.3f' % scores[1])

Epoch 1/5
391/391 [=====] - 6795.2s/step - loss: 1.5665 - accuracy: 0.4357 - val_loss: 1.2536 - val_accuracy: 0.5576
Epoch 2/5
391/391 [=====] - 6785.2s/step - loss: 1.2036 - accuracy: 0.5739 - val_loss: 0.9769 - val_accuracy: 0.6265
Epoch 3/5
391/391 [=====] - 6745.2s/step - loss: 1.0158 - accuracy: 0.6408 - val_loss: 0.9540 - val_accuracy: 0.6616
Epoch 4/5
391/391 [=====] - 6855.2s/step - loss: 0.8619 - accuracy: 0.6981 - val_loss: 0.9144 - val_accuracy: 0.6817
Epoch 5/5
391/391 [=====] - 6815.2s/step - loss: 0.7397 - accuracy: 0.7451 - val_loss: 0.8531 - val_accuracy: 0.7066
313/313 [=====] - 139.626s/step - loss: 0.8531 - accuracy: 0.7066
Loss: 0.853
Accuracy: 0.707

This was very similar to Dropout but a slightly smaller loss. This makes sense since earlier we established that max norm didn't make much of a difference
```

```
In [ ]:
```